Machine Learning: Introduction

Sidharth Baskaran

June 2021

Introduction

- Supervised learning - right answers are given
  - Algorithm is given the correct/expected answer
  - Regression - predict continuous valued output
  - Classification - predict results to a discrete output
    * Can have more than 2 classifications
    * Models may require infinite number of attributes or features, is done with SVM (support vector machine)
- Unsupervised learning - dataset is not classified, a structure must be predicted
  - Example - cluster classification of news articles
  - Approach problems without an idea of the results or knowledge of effect of variables
  - Derived from clustering data based on variable relationships
  - No feedback

Model and Cost Function

- Linear regression algorithm
  - Fitting a line to data - supervised learning
    * Predicting a real-valued output
- Notation
  - m is number of training examples
  - x represents the input variable/features
  - y represents output/target variable
  - (x, y) represents a single training example
  - (x_i, y_i) refers to ith training example

Linear Regression

- Process flow
  - Training set -> learning algorithm -> h
  - h is the hypothesis, function which takes input x and outputs estimated y
    * Maps $x \rightarrow y$
- $h_\theta(x) = \theta_0 + \theta_1 x$ is the cost function
  - θ_i are parameters that correspond to the regression line
    * Choose θ_0, θ_1 so $h_\theta(x)$ is close to y for examples in training data (x, y)
- Minimizing the average of training set residuals

θ_0, θ_1 distance to true values which is minimizing

$$J\left(\theta_0, \theta_1\right) = \frac{1}{2m} \sum_{i=1}^{m} \left(\hat{y}_i - y_i\right)^2 = \frac{1}{2m} \sum_{i=1}^{m} \left(h_\theta\left(x_i\right) - y_i\right)^2$$

- m is the training set size
- Is the squared error cost function - goal is to minimize
- Halving the mean is for convenience as derivative will cancel it
- Hypothesis is a function of x for some fixed θ_1 and $J(\theta_1)$ is a function of θ_1

Contour plots

- 2 parameters θ_0, θ_1 -> 3D plot paraboloid
  - Height is J
- Can find minimum from contour plot
  - Closer to minimum on a contour plot means better fit

Gradient Descent Algorithm

- Outline
  - Want $\min_{\theta_0, \theta_1} J(\theta_0, \theta_1)$
  - Start with some θ_0, θ_1
  - Keep changing θ_0, θ_1 to reduce J until a minimum is reached (for any cost function J)
- Gradient $-\nabla J$ points in direction of steepest *descent*

repeat until convergence {
$$\theta_j := \theta_j - \alpha \frac{\partial}{\partial \theta_j} J\left(\theta_0, \theta_1\right) \quad (\text{ for } j = 0 \text{ and } j = 1)$$
}

- Assignment := is not same as =
  - Can do $a := a + 1$ but not $a = a + 1$
- Simulataneous update must be used

$$\text{temp } 0 := \theta_0 - \alpha \frac{\partial}{\partial \theta_0} J\left(\theta_0, \theta_1\right)$$
$$\text{temp } 1 := \theta_1 - \alpha \frac{\partial}{\partial \theta_1} J\left(\theta_0, \theta_1\right)$$
$$\theta_0 := \text{ temp } 0$$
$$\theta_1 := \text{ temp } 1$$

- $\alpha > 0$ is the learning rate and $\frac{\partial}{\partial \theta_j} J(\theta_0, \theta_1$ is derivative in direction of θ_j
  - A large α means minimum can be overshot
    * Fail to converge -> even diverge
  - Small α means slow descent
- Convergence can occur even with a fixed α since the partial derivative term decreases when minima is approached over time

Gradient Descent for linear regression

- Need to minimize square error cost function

$$j = 0 : \frac{\partial}{\partial \theta_0} J\left(\theta_0, \theta_1\right) = \frac{1}{m} \sum_{i=1}^{m} \left(h_\theta\left(x^{(i)}\right) - y^{(i)}\right)$$
$$j = 1 : \frac{\partial}{\partial \theta_1} J\left(\theta_0, \theta_1\right) = \frac{1}{m} \sum_{i=1}^{m} \left(h_\theta\left(x^{(i)}\right) - y^{(i)}\right) \cdot x^{(1)}$$

- Linear regression always results in a convex (bow) cost function J
  - Will always converge to the global minimum
- Batch gradient descent - each step uses all training examples

Linear Algebra Review

- Dimension of a matrix is row $\times$ col or $\mathbb{R}^{\text{row} \times \text{col}}$
- A_{ij} is ith row and jth column entry of matrix A
- A vector is a $n \times 1$ matrix of dimension n
  - y_i is the ith element in the vector $\vec{y}$
  - Can be 0 or 1-indexed
- Given $h_\theta(x) = \theta_0 + \theta_1 x$, can use $\vec{p} = D \times \vec{\theta}$ where $\vec{p}$ is the predicted regression values vector of dimension 4, D is $n \times 2$ matrix with column $1 = \vec{1}$, and $\vec{\theta}$ is the parameter matrix of dimension 2
  - More computationally efficient
- If $C = A \times B$, the ith column of C is $A \times \vec{B_i}$ where B_i is the ith column of B
- If applying multiple hypotheses to a data set, use a $2 \times n$ matrix where there are n hypotheses and 2 parameters
- No commutative matrix/vector multiplication but associativity works
- Identity matrix: $A \times I = I \times A = A$ but $AB \neq BA$ if $B \neq A \neq I$
- If A is $m \times m$ and has inverse A^{-1}, then $AA^{-1} = A^{-1}A = I$
- A transpose makes the the ith row the ith column and is reversible
  - If A is an $m \times n$ matrix and $B = A^T$, then B is an $n \times m$ matrix and $B_{ij} = A_{ji}$