

Learning with Neural Networks

Sidharth Baskaran

June 2021

Backpropagation Cost Function

- Notation
 - L = total no. of layers in network
 - s_l = no. of units excluding bias in layer l
- Classification types
 - Binary - $y \in \{0, 1\}$
 - Multi-class - $y \in \mathbb{R}^K$ where there are K output units ($h_{\Theta}(x) \in \mathbb{R}^K : K \geq 3$)
- Cost function with regularization
 - Let $(h_{\Theta}(x))_i$ be i th output

$$J(\Theta) = -\frac{1}{m} \sum_{i=1}^m \sum_{k=1}^K \left[y_k^{(i)} \log(h_{\Theta}(x))_k + (1 - y_k^{(i)}) \log(1 - (h_{\Theta}(x))_k) \right] + \frac{\lambda}{2m} \sum_{l=1}^{L-1} \sum_{i=1}^{s_l} \sum_{j=1}^{s_{l+1}} (\theta_{j,i}^{(l)})^2$$

Backpropagation Algorithm

- Need to find $\min_{\Theta} J(\Theta)$
 - Compute $J(\Theta)$ and $\frac{\partial}{\partial \Theta_{ij}^{(l)}}$
- Gradient **forward propagation** computation given 1 example (x, y)
 - $a^{(1)} = x$
 - $z^{(2)} = \Theta^{(1)} a^{(1)}$
 - $a^{(2)} = g(z^{(2)})$ (add $a_0^{(2)}$)
 - etc.
- Intuition - compute $\delta_j^{(l)}$ is error of node j in layer l
 - If output unit is layer i , $\delta_j^{(i)} = a_j^{(i)} - y_j$
 - * $a_j^{(4)} = (h_{\Theta}(x))_j$
 - $\delta^{(i)} = (\Theta^{(i)})^T \delta^{(i+1)} \cdot g'(z^{(i)})$
 - * Can be shown that $g'(z^{(i)}) = a^{(i)} \cdot (1 - a^{(i)})$
 - * No $\delta^{(1)}$ term because input features don't have error
- General training set of m examples $\{(x^{(1)}, y^{(1)}), \dots, (x^{(m)}, y^{(m)})\}$
 - Set $\Delta_{ij}^{(l)} = 0 \forall l, i, j$ (used to compute the partial derivative of $J(\Theta)$)

Algorithm

For $i = 1 \rightarrow m$ ($\leftarrow (x^{(i)}, y^{(i)})$):

- Set $a^{(1)} = x^{(i)}$
- Forward propagation to compute $a^{(l)}$ for $l \in \text{range}(2, 3, \dots, L)$

- Use $y^{(i)}$ to compute $\delta^{(L)} - y^{(i)}$
- Back propagation to compute $\delta^{(L-1)}, \dots, \delta^{(2)}$ (no $\delta^{(1)}$)
- $\Delta_{ij}^{(l)} := \Delta_{ij}^{(l)} + a_j^{(l)} \delta_i^{(l+1)} \longleftrightarrow \Delta^{(l)} := \Delta^{(l)} + \delta^{(l+1)}(a^{(l)})^T$

Then

$$D_{ij}^{(l)} := \frac{1}{m} (\Delta_{ij}^{(l)} + \lambda \Theta_{ij}^{(l)}) | j \neq 0$$

$$D_{ij}^{(l)} := \frac{1}{m} \Delta_{ij}^{(l)} | j = 0$$

- Can then show that $\boxed{\frac{\partial}{\partial \Theta_{ij}^{(l)}} J(\Theta)} = D_{ij}^{(l)}$
- Use this partial derivative in an algorithm like gradient descent or `fminunc`

Backpropagation Intuition

- Forward propagation - L to R
 - Inputs from data set fed into 1st layer
 - Sigmoid applied to z values to get activation values in each layer
 - Weights (arrow/lines) \times originating activation values (a dot product) \rightarrow new z -value to which sigmoid is applied
- Backpropagation - R to L
 - Let i th training example have a cost $\text{cost}(i) = y^{(i)} \log h_{\Theta}(x^{(i)}) + (1 - y^{(i)}) \log h_{\Theta}(x^{(i)})$
 - * Measure of how well the network is performing on example i
 - * Ignore regularization, so $\lambda = 0$
 - $\delta_j^{(l)}$ is error of cost for $a_j^{(l)}$
 - Formally, $\delta_j^{(l)} = \frac{\partial}{\partial z_j^{(l)}} \text{cost}(i)$ for $j \geq 0$ so **delta values are derivative of the cost function**
 - * Steeper slope means more incorrect, so is defined as δ
- Backpropagation in the sense that an error of a unit in previous layer is weighted sum (using Θ) of errors in current layer that are edged to the unit in question
 - Opposite of forward-propagation

Implementation note - unrolling parameters

- Process of unrolling parameters from matrices \rightarrow vectors

Optimization routine:

```
function [jval, gradient] = costFunction(theta)
...
optTheta = fminunc(@costFunction, initialTheta, options)
```

Neural network with ex. $L = 4$: $\Theta^{(1)}, \Theta^{(2)}, \Theta^{(3)} \rightarrow$ matrices `Theta1`, `Theta2`, `Theta3` $D^{(1)}, D^{(2)}, D^{(3)} \rightarrow$ matrices `D1`, `D2`, `D3`

Unroll into vectors:

```
thetaVec = [ Theta1(:); Theta2(:); Theta3(:) ];
gradientVec = [ D1(:); D2(:); D3(:) ];
```

Getting back to matrices:

```
Theta1 = reshape(thetaVec(1:110), 10, 11);
Theta2 = reshape(thetaVec(111:220), 10, 11);
Theta3 = reshape(thetaVec(221:231), 1, 11);
```

Can then do

```
function [jval, gradientVec] = costFunction(thetaVec);
```

Use unrolling and forward/back propagation to get $D^{(i)}$ and $J(\Theta)$

- Example NN
 - $s_1 = 10, s_2 = 10, s_3 = 1$
 - $\Theta^{(1)} \in \mathbb{R}^{10 \times 11}, \Theta^{(2)} \in \mathbb{R}^{10 \times 11}, \Theta^{(3)} \in \mathbb{R}^{1 \times 11}$
 - $D^{(1)} \in \mathbb{R}^{10 \times 11}, D^{(2)} \in \mathbb{R}^{10 \times 11}, D^{(3)} \in \mathbb{R}^{1 \times 11}$

Gradient Checking

- Allows to check implementation of backprop
- Numerical estimation of gradients
 - $\frac{d}{d\theta} J(\theta) \approx \frac{J(\theta+\epsilon) - J(\theta-\epsilon)}{2\epsilon}$ where $\epsilon = 10^{-4}$
 - Implement this as some `gradApprox` in Octave
- If $\theta = [\theta_1, \dots, \theta_n] \in \mathbb{R}^n$, then the n th partial derivative is $\frac{\partial}{\partial \theta_n} J(\theta) \approx \frac{J(\theta_1, \dots, \theta_n + \epsilon) - J(\theta_1, \dots, \theta_n - \epsilon)}{2\epsilon}$

Algorithm:

```
for i = 1:n,  
    % setting theta +/- epsilon  
    thetaPlus = theta;  
    thetaPlus(i) = thetaPlus(i) + EPSILON;  
    thetaMinus = theta;  
    thetaMinus(i) = thetaMinus(i) - EPSILON;  
    % grad approximation  
    gradApprox(i) = (J(thetaPlus) - J(thetaMinus))/(2 * EPSILON)  
end  
% example check  
gradApprox - DVec <= 10^(-4) % precision desired
```

- Implementation notes
 - Implement backprop to compute `DVec` (unrolled D matrices)
 - Implement `gradApprox` through algorithm and check with gradient
 - Turn off checking and use backprop for learning
 - Make sure to disable gradient check \rightarrow slows down code significantly due to numerical gradient computations on every optimization iteration - **reason for using backprop**

Random Initialization

- Pick initial value for Θ
 - Can not initialize to 0, makes activations, errors, and partials equivalent from 1st to 2nd layer
 - All hidden features compute same function of input \rightarrow redundant
- Initialize each $\Theta_{ij}^{(l)}$ to random value $\in [-\epsilon, \epsilon]$
- Values should be close to 0 but not 0
- Need to break symmetry

```
Theta1 = rand(10,11) * (2 * INIT_EPSILON) - INIT_EPSILON;  
Theta2 = rand(1,11) * (2 * INIT_EPSILON) - INIT_EPSILON;
```

Overall Implementation Process

- Network architecture - connectivity between neurons

- No. of input units - dimension of features $x^{(i)}$
- No. of output units - number of classes
 - * $y \in \mathbb{R}^K$ where there are K classes
- Default - 1 hidden layer and if ≥ 1 , equal number of units in each layer
 - * More is better but balance with cost function
- Training neural network steps
 - Randomly initialize weights
 - Implement forward propagation to get $h_{\Theta}(x^{(i)})$ for any $x^{(i)}$
 - Compute $J(\Theta)$
 - Implement backprop to get partial derivatives of $J(\Theta)$
 - Use gradient checking to compare backpropagation partial with numerical estimate of $\nabla J(\Theta)$, then disable check
 - Use optimization routine to minimize $J(\Theta)$

`for i = 1:m,`

`Perform forward propagation and backpropagation using example (x(i),y(i))`

`(Get activations a(l) and delta terms d(l) for l = 2,...,L`

- For NN $\rightarrow J(\Theta)$ is **non-convex** so has local extrema
 - Ideally want $h_{\Theta}(x^{(i)}) \approx y^{(i)}$