Advice for Applying Machine Learning

Sidharth Baskaran

June 2021

Improving algorithm

- If there are large errors in predictions
  - More training examples
  - Smaller set of features
  - Additional features
  - Add polynomial features (i.e. x_1x_2, x_1^2, x_2^2)
  - Decrease or increase λ
- Machine learning diagnostic
  - Test run to understand performance of algorithm

Evaluating a Hypothesis

- Low training error does not mean good $h_\theta(x)$
- Hard to plot hypothesis as $n \to$ large
- Split data into 2 parts $\to$ test and training
- Training/testing procedure for linear regression
  - Leanr parameter θ from training data (minimize training error $J(\theta)$
  - Compute test error
    * $J_{\text{test}}(\theta) = \frac{1}{2m_{\text{test}}} \sum_{i=1}^{m_{\text{test}}} (h_\theta(x_{\text{test}}^{(i)}) - y_{\text{test}}^{(i)})^2$
    * Change appropriately for logistic regression
  - Misclassificaiton error (0/1 misclassificaiton error)
    * Test error $= \frac{1}{m_{\text{test}}} \sum_{i=1}^{m_{\text{test}}} \text{err}(h_\theta(x_{\text{test}}^{(i)}), y_{\text{test}}^{(i)})$
      · Is proportion of misclassified test data

$$\text{err}(h_\theta(x), y) = \begin{cases} 1 \; y = 0 \text{ if } h_\theta(x) \geq 0.5 \textbf{ or } y = 1 \text{ if } h_\theta(x) < 0.5 \\ 0 \text{ otherwise} \end{cases}$$

Model Selection and training/validation/test sets

- Error of parameters as measured to fit to a set of data < than actual generalization error
- Let there be a list of hypotheses of varying polynomial degree d and $i \in \text{range}(1, d)$
  - Calculate $J_{\text{test}}(\theta^{(i)})$ in range
  - Lowest J_{test} means best model $\to$ but not fair estimate of generalization $\to$ optimistic choice
- Split data set into 3 parts
  - E.g. 60% training, 20% cross-validation (CV), 20% test
  - Determine $J_{\text{train}}, J_{\text{CV}}, J_{\text{test}}$
  - In list of polynomial models - find lowest J_{CV} from CV set
  - Estimate generalization error from polynomial indexed d with lowest CV error

Bias vs. Variance

- Increase degree d of polynomial $\rightarrow$ decreases training error of polynomial
  - Test error very similar
- Cross-validation error initially decreases then increases as overfitting begins to occur
- **High bias** problem $\rightarrow$ low d and a high error
  - High $J_{\text{train}}(\theta)$ and $J_{\text{CV}}(\theta) \approx J_{\text{train}}(\theta)$
- **High variance** problem $\rightarrow$ high d and a high error
  - $J_{\text{train}}(\theta)$ is low
  - $J_{\text{CV}}(\theta) \gg J_{\text{train}}(\theta)$

Regularization and bias/variance

- High λ means penalized parameters θ so high bias $\rightarrow$ underfit
- Intermediate $\lambda \rightarrow$ optimal
- Small $\lambda \rightarrow$ high variance and overfit
- Choosing regularization parameter λ
  - $J_{\text{train}}(\theta), J_{\text{CV}}(\theta), J_{\text{test}}(\theta)$ all do not have regularization term but $J(\theta)$ does
  - Try a range of λ, e.g. in multiples of 2 and calculate the corresponding $J_{\text{CV}}(\theta)$
    - Pick choice with lowest CV error
  - Then apply to $J_{\text{test}}(\theta)$ to check for good generalization
- Can then plot $J_{\text{train}}(\theta)$ and $J_{\text{CV}}(\theta)$ as a function of λ
  - $J_{\text{train}}(\theta)$ will increase
  - $J_{\text{CV}}(\theta)$ will be upward parabolic

Learning Curves

- Plot either training or CV error
- A small training set size m means virtually no error
  - As m increwases, avg. training error of hypotheses increases
- CV error is high (low generalization) $\rightarrow$ small m
  - Tend to decrease with m
- More training data does not help case of high bias
- High variance problem
  - CV error will decrease with higher m
- As m increases, $J_{\text{CV}}(\theta) - J_{\text{train}}(\theta)$ and both curves approach each other

High bias

Value of m	$J_{\text{train}}(\theta)$	$J_{\text{CV}}(\theta)$
Low	Low	High
High	High	Low

High variance

Value of m	$J_{\text{train}}(\theta)$	$J_{\text{CV}}(\theta)$
Low	Low	High
High	Increases	Decreases

Debugging learning algorithm

- Choices
  - More training examples (m) $\rightarrow$ fixes high variance
  - Smaller sets of features $\rightarrow$ fixes high variance
  - Additional features $\rightarrow$ fixed high bias
  - Adding polynomial features $\rightarrow$ fixes high bias
  - Decreasing λ $\rightarrow$ fixes high bias
  - Increasing λ $\rightarrow$ fixed high variance
- Small neural network $\rightarrow$ prone to underfitting due to less parameters, computationally cheaper
- Large neural network $\rightarrow$ prone to overfitting and computationally expensive
  - Can address with λ