

# Advice for Applying Machine Learning

Sidharth Baskaran

June 2021

## Improving algorithm

- If there are large errors in predictions
  - More training examples
  - Smaller set of features
  - Additional features
  - Add polynomial features (i.e.  $x_1x_2, x_1^2, x_2^2$ )
  - Decrease or increase  $\lambda$
- Machine learning diagnostic
  - Test run to understand performance of algorithm

## Evaluating a Hypothesis

- Low training error does not mean good  $h_\theta(x)$
- Hard to plot hypothesis as  $n \rightarrow$  large
- Split data into 2 parts  $\rightarrow$  test and training
- Training/testing procedure for linear regression
  - Learn parameter  $\theta$  from training data (minimize training error  $J(\theta)$ )
  - Compute test error
    - \*  $J_{\text{test}}(\theta) = \frac{1}{2m_{\text{test}}} \sum_{i=1}^{m_{\text{test}}} (h_\theta(x_{\text{test}}^{(i)}) - y_{\text{test}}^{(i)})^2$
    - \* Change appropriately for logistic regression
  - Misclassification error (0/1 misclassification error)
    - \* Test error =  $\frac{1}{m_{\text{test}}} \sum_{i=1}^{m_{\text{test}}} \text{err}(h_\theta(x_{\text{test}}^{(i)}), y_{\text{test}}^{(i)})$
    - Is proportion of misclassified test data

$$\text{err}(h_\theta(x), y) = \begin{cases} 1 & y = 0 \text{ if } h_\theta(x) \geq 0.5 \text{ or } y = 1 \text{ if } h_\theta(x) < 0.5 \\ 0 & \text{otherwise} \end{cases}$$

## Model Selection and training/validation/test sets

- Error of parameters as measured to fit to a set of data < than actual generalization error
- Let there be a list of hypotheses of varying polynomial degree  $d$  and  $i \in \text{range}(1, d)$ 
  - Calculate  $J_{\text{test}}(\theta^{(i)})$  in range
  - Lowest  $J_{\text{test}}$  means best model  $\rightarrow$  but not fair estimate of generalization  $\rightarrow$  optimistic choice
- Split data set into 3 parts
  - E.g. 60% training, 20% cross-validation (CV), 20% test
  - Determine  $J_{\text{train}}, J_{\text{CV}}, J_{\text{test}}$
  - In list of polynomial models - find lowest  $J_{\text{CV}}$  from CV set
  - Estimate generalization error from polynomial indexed  $d$  with lowest CV error

## Bias vs. Variance

- Increase degree  $d$  of polynomial  $\rightarrow$  decreases training error of polynomial
  - Test error very similar
- Cross-validation error initially decreases then increases as overfitting begins to occur
- **High bias** problem  $\rightarrow$  low  $d$  and a high error
  - High  $J_{\text{train}}(\theta)$  and  $J_{\text{CV}}(\theta) \approx J_{\text{train}}(\theta)$
- **High variance** problem  $\rightarrow$  high  $d$  and a high error
  - $J_{\text{train}}(\theta)$  is low
  - $J_{\text{CV}}(\theta) \gg J_{\text{train}}(\theta)$

## Regularization and bias/variance

- High  $\lambda$  means penalized parameters  $\theta$  so high bias  $\rightarrow$  underfit
- Intermediate  $\lambda \rightarrow$  optimal
- Small  $\lambda \rightarrow$  high variance and overfit
- Choosing regularization parameter  $\lambda$ 
  - $J_{\text{train}}(\theta)$ ,  $J_{\text{CV}}(\theta)$ ,  $J_{\text{test}}(\theta)$  all do not have regularization term but  $J(\theta)$  does
  - Try a range of  $\lambda$ , e.g. in multiples of 2 and calculate the corresponding  $J_{\text{CV}}(\theta)$ 
    - \* Pick choice with lowest CV error
  - Then apply to  $J_{\text{test}}(\theta)$  to check for good generalization
- Can then plot  $J_{\text{train}}(\theta)$  and  $J_{\text{CV}}(\theta)$  as a function of  $\lambda$ 
  - $J_{\text{train}}(\theta)$  will increase
  - $J_{\text{CV}}(\theta)$  will be upward parabolic

## Learning Curves

- Plot either training or CV error
- A small training set size  $m$  means virtually no error
  - As  $m$  increases, avg. training error of hypotheses increases
- CV error is high (low generalization)  $\rightarrow$  small  $m$ 
  - Tend to decrease with  $m$
- More training data does not help case of high bias
- High variance problem
  - CV error will decrease with higher  $m$
- As  $m$  increases,  $J_{\text{CV}}(\theta) - J_{\text{train}}(\theta)$  and both curves approach each other

### High bias

Value of $m$	$J_{\text{train}}(\theta)$	$J_{\text{CV}}(\theta)$
Low	Low	High
High	High	Low

### High variance

Value of $m$	$J_{\text{train}}(\theta)$	$J_{\text{CV}}(\theta)$
Low	Low	High
High	Increases	Decreases

## Debugging learning algorithm

- Choices
  - More training examples ( $m$ )  $\rightarrow$  fixes high variance
  - Smaller sets of features  $\rightarrow$  fixes high variance
  - Additional features  $\rightarrow$  fixed high bias
  - Adding polynomial features  $\rightarrow$  fixes high bias
  - Decreasing  $\lambda \rightarrow$  fixes high bias
  - Increasing  $\lambda \rightarrow$  fixed high variance
- Small neural network  $\rightarrow$  prone to underfitting due to less parameters, computationally cheaper
- Large neural network  $\rightarrow$  prone to overfitting and computationally expensive
  - Can address with  $\lambda$

## Machine Learning System Design

- Supervised learning  $\rightarrow x =$  features of email,  $y \in \{0, 1\}$ , can choose 100 words indicative of spam/not for features
- Can encode an email into a feature vector
  - In practice  $\rightarrow$  take most frequently occurring  $n$  words in training set
- Could develop features to reduce errors, e.g. misspelling detection
  - Equal consideration of all options  $\rightarrow$  cannot tell which will work best

## Error Analysis

- Start with **simple** algorithm to test on CV data
- Plot learning curves to decide if more data, features, etc.
- Error analysis  $\rightarrow$  manual examination of examples where errors occurred
  - Look for systematic error trend
  - Use evidence to guide decision-making not guesswork
- Numerical evaluation
  - Treating stem of word = to variants of word
  - Can use stemming software
  - Naturally  $\rightarrow$  CV error  $J_{CV}(\theta)$  of algorithm with/without stemming and choose best options
    - \* Do not  $J_{train}(\theta)$  to allow for generalization

## Error metrics for skewed classes

- Precision/recall
  - **Precision**  $\rightarrow$  Of all predictions  $y = 1$ , fraction that actually correspond to  $y = 1$ 
    - \*  $\text{Precision} = \frac{\text{True pos.}}{\text{Predicted pos.}} = \frac{\text{True pos.}}{\text{True pos.} + \text{False pos.}}$
  - **Recall**  $\rightarrow$  Of all actual cases  $y = 1$ , what fraction was correctly detected by algorithm
    - \*  $\text{Recall} = \frac{\text{True pos.}}{\text{Actual pos.}} = \frac{\text{True pos.}}{\text{True pos.} + \text{False neg.}}$
- Tradeoff of precision/recall
  - If trying to have high confidence  $\rightarrow$  high precision, low recall
  - If trying to minimize false negatives  $\rightarrow$  high recall, low precision
- In general  $\rightarrow$  predict 1 if  $h_{\theta}(x) \geq$  threshold
- $F_1$  score  $\rightarrow$  comparing precision/recall numbers
  - Can calculate average  $\frac{P+R}{2}$  but susceptible to high or low recall/precision weighted
  - $F_1$  or  $F$ -score better  $\rightarrow 2 \frac{PR}{P+R}$ 
    - \* Gives more weight to  $\min(P, R)$

## Data for Machine Learning

- Large data rationale
  - Assume feature  $x \in \mathbb{R}^{n+1}$  has enough information to predict  $y$  accurately
  - Can ask if given input  $x$ , can human expert confidently predict  $y$
  - A learning algorithm with many parameters or NN with many hidden layers
    - \*  $J_{\text{train}}(\theta)$  very small
  - Very large training set  $\rightarrow$  unlikely to overfit
    - \*  $J_{\text{train}}(\theta) \approx J_{\text{test}}(\theta)$